DEVELOPING AN EMOTION DETECTION SYSTEM IN SPEECH FOR ENHANCED REAL-TIME APPLICATIONS

A PROJECT REPORT

Submitted in Partial Fulfillment for the degree of
Bachelor of Technology
In
ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING
From
MAULANA ABUL KALAM AZAD UNIVERSITY OF TECHNOLOGY

MAULANA ABUL KALAM AZAD
UNIVERSITY OF TECHNOLOGY,
WEST BENGAL

Submitted by

SHOUNAK SARKAR	**11600921047**
PRIYAM PAL	**11600921032**
AMRIT BAG	**11600921005**
SAPTARSHI PARUI	**11600921042**
SAGAR DAS	**11600921036**
SAYAK SINHA ROY	**11600921046**

Under the guidance of
MR. PRIYANATH MAHANTI
Assistant professor of IT Department

DEPARTMENT OF INFORMATION TECNOLOGY
MCKV Institute of Engineering
243, G.T. Road (North), Liluah, Howrah-711204
(NAAC Accredited Grade- A Autonomous Institute affiliated to
Maulana Abul Kalam Azad University of Technology, West Bengal)

May, 2025

DEPARTMENT OF INFORMATION TECHNOLOGY
MCKV Institute of Engineering
An Autonomous Institute Affiliated to
Maulana Abul Kalam Azad University of Technology, West Bengal

CERTIFICATE

This is to certify that the project report entitled ***"DEVELOPING AN EMOTION DETECTION SYSTEM IN SPEECH FOR ENHANCED REAL-TIME APPLICATIONS"*** *submitted by "***SHOUNAK SARKAR, PRIYAM PAL, AMRIT BAG, SAPTARSHI PARUI, SAGAR DAS, SAYAK SINHA ROY*** *for* 8^{th} *semester examination has been prepared for the partial fulfillment of the degree of Bachelor of Technology in Artificial Intelligence and Machine Learning, awarded by the Maulana Abul Kalam Azad University of Technology, West Bengal. They have carried out the project work under my supervision.*

-- --

HEAD OF THE DEPARTMENT

Mr. Sachin Balo

Dept. of Information Technology
MCKV Institute of Engineering
243, G.T. Road (North), Liluah
Howrah – 711204

SUPERVISOR

Mr. Priyanath Mahanti

Dept. of Information Technology
MCKV Institute of Engineering
243, G.T. Road (North), Liluah
Howrah – 711204

DEPARTMENT OF INFORMATION TECHNOLOGY
MCKV Institute of Engineering

An Autonomous Institute Affiliated to

Maulana Abul Kalam Azad University of Technology, West Bengal

CERTIFICATE OF APPROVAL

(B.Tech Degree in Artificial Intelligence and Machine Learning)

This project report is hereby approved as a creditable study of an engineering subject carried out and presented in a manner satisfactory to warrant its acceptance as a prerequisite to the degree for which it has been submitted. It is to be understood that by this approval, the undersigned do not necessarily endorse or approve any statement made, opinion expressed and conclusion drawn therein but approve the project report only for the purpose for which it has been submitted

SIGNATURE OF EXAMINATION COMMITTEE MEMBERS FOR EVALUATION OF PROJECT REPORT	1. Internal Examiner 2. Internal Examiner 3. Internal Examiner 4. Internal Examiner 5. External Examiner

Date of evaluation:

ACKNOWLEDGEMENT

It is always a pleasure to express my sincere thanks to those who have supported us throughout this work. we hereby acknowledge the guidance and support of **Mr. Priyanath Mahanti, Assistant professor of IT Department,** whose expertise and encouragement were invaluable to the completion of this project. we are also grateful to **Mr. Sachin Balo, Head of the Department (Information Technology) and other supporting faculty and staff members** for providing the necessary support and infrastructure.

_ _

Shounak Sarkar **Priyam Pal**

Amrit Bag **Saptarshi Parui**

Sagar Das **Sayak Sinha Roy**

ABSTRACT

Emotion detection from speech is increasingly essential for enhancing real-time applications in fields such as healthcare, customer service, and human-computer interaction. Traditional emotion recognition approaches struggle with challenges like data variability, real-time constraints, and the complexity of contextual understanding. This research presents an advanced emotion detection system that integrates Convolutional Neural Networks (CNN), Recurrent Neural Networks (RNN), and Transformer networks, alongside other state-of-the-art algorithms, to address these limitations.

Our system uses a multi-stage approach to leverage the strengths of different models. CNNs are employed for spectral feature extraction, while RNNs capture the temporal dynamics in speech data. Transformer networks add the ability to model long-range dependencies, enhancing the system's understanding of complex and subtle emotions. We incorporate a comprehensive feature extraction pipeline, including Mel-Frequency Cepstral Coefficients (MFCC), pitch, and prosodic features, to maximize emotional context and model robustness.

Extensive evaluations are conducted on publicly available datasets to benchmark the model's real-time performance, classification accuracy, and robustness across various environmental conditions. The results highlight that this hybrid architecture outperforms traditional models, offering a balance between accuracy, latency, and computational efficiency for real-time deployment. Our findings suggest that the Transformer-based approach provides an additional advantage in cross-cultural and contextually complex emotion scenarios, addressing prior gaps in real-world applicability.

This research provides a robust framework for real-time emotion detection, offering valuable implications for practical applications, including virtual assistants, real-time mental health monitoring, and sentiment analysis in customer support. Future work aims to incorporate multi-modal data sources, explore cross-linguistic adaptability, and address ethical considerations in emotion detection technology.

Keywords: Emotion Detection, Speech Recognition, Real-Time Applications, Deep Learning, CNN, RNN, Transformer Network, Feature Extraction, Latency Reduction, MFCC.

Contents

LIST OF TABLES

LIST OF FIGURES

1. INTRODUCTION

1.1 Background and Motivation

Emotion detection from speech has gained significant attention as technology moves towards more personalized, human-centered applications. In healthcare, emotion detection can be transformative, aiding in the monitoring and treatment of mental health issues such as anxiety, depression, and stress. Real-time emotion detection systems enable clinicians to track emotional changes, providing insights into patients' psychological states over time. This can be particularly valuable in telemedicine and digital mental health platforms, where immediate emotional assessment could facilitate more responsive care.

In human-computer interaction (HCI), emotion detection enhances the user experience by making technology more responsive to human emotions. By understanding emotional cues, virtual assistants, customer service bots, and interactive systems can offer responses that feel more empathetic and contextually appropriate, increasing user satisfaction and engagement. Emotion-aware systems could, for example, detect frustration in a customer service scenario, triggering an escalation to a human representative or adjusting the tone of the conversation.

Beyond healthcare and HCI, real-time emotion detection holds potential in education, where adaptive learning platforms can use emotional feedback to adjust content delivery, ensuring students remain engaged and reducing frustration. Similarly, in entertainment and social media, emotion-aware systems could enrich user interactions, making platforms more engaging and relevant.

Despite its potential, developing an accurate and real-time emotion detection system remains challenging. Current systems often struggle with factors like data diversity, cross-cultural variations, and latency in real-time applications. These challenges motivate our research to create a robust, adaptable, and scalable emotion detection system, capable of operating effectively across various real-world scenarios.

1.2 Problem Statement

This research aims to develop a real-time emotion detection system from speech, addressing key challenges like data variability, cultural and contextual adaptation, and latency in processing. By leveraging CNN, RNN, and Transformer architectures, the system seeks to improve emotion recognition accuracy and speed, making it suitable for real-world applications in healthcare, human-computer interaction, and adaptive user interfaces.

1.3 Objectives

The primary objective of this research is to develop a robust and efficient emotion detection system that operates in real-time, offering improvements over current systems in terms of accuracy, adaptability, and processing speed. The specific objectives include:

Improving Emotion Recognition Accuracy: To enhance the accuracy of emotion detection by integrating deep learning models such as Convolutional Neural Networks (CNN), Recurrent Neural Networks (RNN), and Transformer networks, each addressing different aspects of speech data like spectral, temporal, and contextual features. This hybrid approach aims to capture a wider range of emotional expressions and improve the system's ability to detect nuanced emotional states.

Real-Time Processing Optimization: To minimize latency and ensure efficient real-time emotion recognition in diverse environments, making it suitable for applications such as mental health monitoring, customer service, and human-computer interaction. This includes optimizing the system's response time without compromising accuracy.

Cross-Cultural and Contextual Adaptability: To design a system that adapts to variations in speech patterns caused by different languages, cultures, and individual characteristics like age and gender. This would ensure broader applicability in diverse real-world settings, from global customer service platforms to healthcare systems.

Comprehensive Feature Extraction: To investigate and implement various feature extraction techniques, such as Mel-Frequency Cepstral Coefficients (MFCC), pitch, and prosody, to capture the most relevant speech characteristics for emotion detection. This includes analyzing how different features impact the performance of various algorithms.

Evaluation and Benchmarking: To rigorously evaluate the proposed system using publicly available emotion-labeled datasets and compare it against state-of-the-art emotion detection systems. This objective focuses on demonstrating the superior performance of the hybrid model in terms of accuracy, robustness, and processing efficiency.

By achieving these objectives, this research aims to contribute to the advancement of emotion detection systems that can be seamlessly integrated into real-time applications across various fields, from healthcare and customer service to interactive technologies.

2. LITERATURE REVIEW

2.1 Historical Development

The historical development of Speech Emotion Recognition (SER) has evolved significantly over the past few decades, from early statistical models to the latest deep learning techniques.

Initially, SER systems relied heavily on machine learning algorithms like Gaussian Mixture Models (GMMs) and Support Vector Machines (SVMs) for classifying emotions based on acoustic features such as pitch, tone, and tempo. These models typically used handcrafted features extracted from the speech signals to predict emotional states. Over time, researchers incorporated more complex features, including voice quality and prosody, to improve accuracy.

In the 2000s, deep learning approaches began gaining traction in SER. Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), especially Long Short-Term Memory (LSTM) networks, allowed for better feature extraction and the modeling of temporal dependencies, which significantly improved the performance of emotion detection systems. These models were able to process raw speech data without the need for explicit feature engineering.

More recently, transformer-based architectures have emerged, offering advantages in capturing long-range dependencies and contextual information. These models use attention mechanisms, which help focus on the most relevant parts of the speech signal for emotion recognition. Transformer networks have shown promising results, particularly in handling the complexities of natural language and emotional expressions.

Year	Development/Key Event	Description
2000s	Introduction of Neural Networks	Neural networks were introduced as a more powerful method for emotion recognition from speech, offering better classification accuracy than traditional machine learning techniques.
2010	First Use of Convolutional Neural Networks (CNNs) for SER	CNNs were applied to audio feature extraction, significantly improving performance in recognizing emotional speech from large datasets.
2015	Recurrent Neural Networks (RNNs) and LSTM	RNNs, especially Long Short-Term Memory (LSTM) networks, were introduced to address the time-dependent nature of speech data, improving real-time emotion detection.
2017	Introduction of Transformer Networks	Transformer-based models, which use attention mechanisms, were explored for SER, bringing a breakthrough in capturing long-range dependencies in speech data.
2020	Multi-Modal Emotion Recognition	Researchers began integrating speech with visual and physiological signals to create more robust emotion recognition systems, leading to advancements in multi-modal emotion recognition.
2022	Real-Time SER Systems	Focus shifted to optimizing emotion recognition models for real-time applications, including healthcare, customer support, and human-computer interaction, with an emphasis on edge computing and low-latency processing.

Table – 1: Milestones in Speech Emotion Recognition

2.2 Current Techniques

Current techniques in Speech Emotion Recognition (SER) predominantly rely on deep learning models like Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), and Transformer networks. CNNs are used for capturing spatial features from speech spectrograms, while RNNs model the temporal dependencies inherent in speech. Transformer networks, with their attention mechanisms, have proven effective in processing long-range dependencies and contextual information in speech data. Additionally, multimodal approaches, which combine speech with other signals like facial expressions, have gained traction to enhance emotion detection accuracy. Recent advancements also include using pre-trained models and transfer learning for improved performance on smaller, domain-specific datasets.

Technique	Model Type	Key Features	Pros	Cons
Convolutional Neural Networks (CNNs)	Deep Learning (Supervised)	Uses spatial hierarchies in features, operates on raw audio spectrograms	Excellent for feature extraction from spectrograms, high accuracy	Requires large amounts of labeled data, slow training
Recurrent Neural Networks (RNNs)	Deep Learning (Supervised)	Captures temporal dependencies in speech data	Effective for sequential data, handles time dependencies well	Prone to vanishing gradient problem, slow training
Long Short-Term Memory (LSTM)	Deep Learning (Supervised)	An advanced RNN, overcomes vanishing gradient problem	Well-suited for emotion recognition in speech with long-range dependencies	Requires significant computational power, training time
Transformer Networks	Deep Learning (Supervised)	Uses self-attention mechanism, highly parallelizable	Faster training due to parallelization, captures long-range dependencies well	Requires large datasets, complex to train and fine-tune
Support Vector Machines (SVM)	Machine Learning (Supervised)	Classifier based on maximum margin, works well with smaller datasets	Effective for small to medium-sized datasets, robust to overfitting	May not perform as well on large, high-dimensional data
Random Forests (RF)	Machine Learning (Supervised)	Ensemble learning method using decision trees	Can handle large datasets with various types of features	Not ideal for real-time applications, may overfit with noisy data

Table – 2: Comparison of Popular SER Techniques

2.3 Unsolved Challenges

Despite significant advancements in Speech Emotion Recognition (SER), several challenges remain unsolved, which contrasts with the approach taken in this research paper. One major issue is data variability, as existing systems struggle to adapt to different accents, languages, and individual characteristics such as age and gender. Additionally, environmental noise remains a critical challenge, as background sounds can significantly degrade emotion detection accuracy.

Another key challenge is real-time processing. While deep learning models like CNNs, RNNs, and Transformers have shown great potential, their computational demands often result in high latency, limiting their application in real-time scenarios

. Moreover, most existing models fail to effectively capture subtle emotional cues in speech, particularly in complex emotional states such as sarcasm or mixed emotions, which this research aims to address with a hybrid approach.

Lastly, generalization across domains is a persistent problem. Current models trained on specific datasets often fail to generalize well when applied to different contexts or cultures, a gap that this research intends to bridge through a more adaptable and robust system.

3. METHODOLOGY

3.1 System Architecture

3.1.1 Overview of System Components

The architecture for a real-time Speech Emotion Recognition (SER) system is composed of the following components, each serving a specific role to ensure efficient and accurate emotion detection:

1. Data Collection & Preprocessing

Purpose: Captures and cleans raw speech data.

Process: Audio input is collected and then processed to remove noise and non-speech elements. Features like Mel-Frequency Cepstral Coefficients (MFCCs), pitch, and intensity are extracted, which are crucial for identifying emotions.

2. Feature Extraction & Fusion

Purpose: Converts raw audio into a structured feature set.

Process: Features representing spectral, prosodic, and temporal aspects are extracted. These are combined in a fusion layer to provide a comprehensive input that captures both static and dynamic emotional cues.

3. Emotion Detection Models

CNNs: Analyze visual representations of audio (spectrograms) for spatial patterns related to emotion.

RNNs/LSTMs: Capture sequential dependencies within speech, useful for modeling time-evolving emotional states.

Transformer Networks: Use attention mechanisms to handle long-range dependencies, enhancing detection of complex emotions.

Ensemble Layer: Fuses model outputs to leverage each model's strengths, improving classification accuracy.

3.1.2 Data Flow and Pipeline Design for Real-Time Applications

To ensure efficient, real-time emotion detection, the data flow and pipeline are designed to handle incoming speech data continuously, from initial capture to final classification, while minimizing latency. Here’s a breakdown of the pipeline:

Input Acquisition: Captures real-time audio data through microphones or streaming services, ensuring low-latency input to avoid delays in processing.

Preprocessing: Processes incoming audio to reduce noise, normalize volume, and extract key features like pitch and MFCCs. This step prepares audio data for efficient model ingestion while retaining emotional cues.

Feature Extraction: Converts processed audio into spectrograms and extracts temporal features. This enables both CNNs and RNNs to effectively handle spatial and sequential aspects of the audio signal.

Model Inference: Sends features to a hybrid model (CNN, RNN, Transformer) pipeline. Each model component processes the data simultaneously or in parallel to optimize latency, providing robust outputs for ensemble fusion.

Ensemble Fusion: Combines predictions from multiple models to improve accuracy and stability of emotion detection. This aggregated result ensures that misclassifications from individual models are minimized.

Real-Time Classification: The ensemble output is passed through a softmax layer to determine the dominant emotion in near real-time. This step is optimized to ensure responses within milliseconds.

Post-Processing: Applies smoothing techniques to stabilize continuous emotion tracking, reducing rapid shifts in detection for a more natural, real-time output.

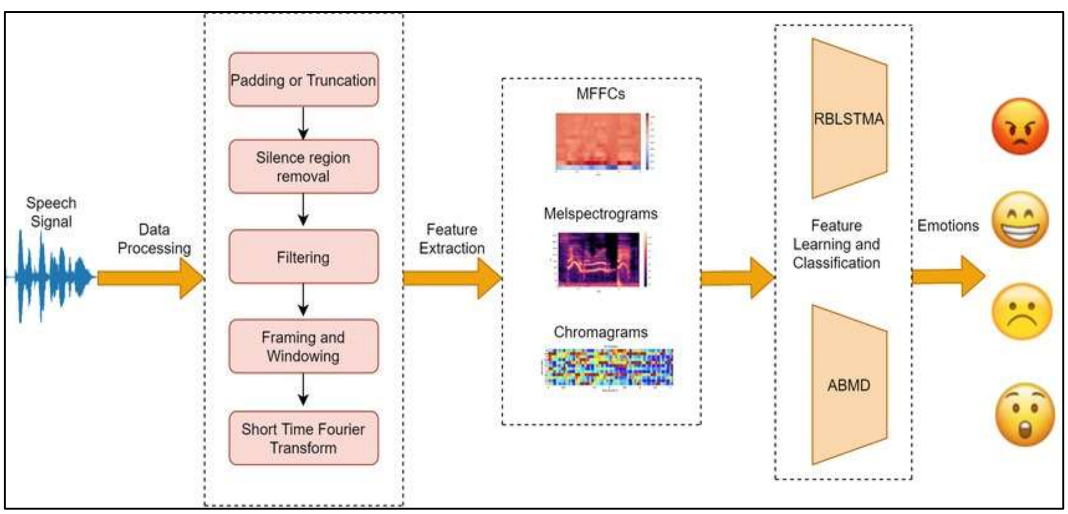

Figure – 1: The Speech Emotion Recognition (SER) system overview.

3.2 Data Collection and Preprocessing

3.2.1 Dataset Selection Criteria and Sources

For this research, the RAVDESS (Ryerson Audio-Visual Database of Emotional Speech and Song) dataset was selected due to its high-quality, validated emotional speech samples, which are essential for building reliable emotion detection models. RAVDESS includes a diverse range of emotions such as happiness, sadness, anger, and calmness, captured in both audio and visual formats, which enhances the dataset's versatility for emotion classification. It's clear labeling and structured format make it suitable for training deep learning models on various emotional states, aligning well with our objective of real-time, accurate emotion detection.

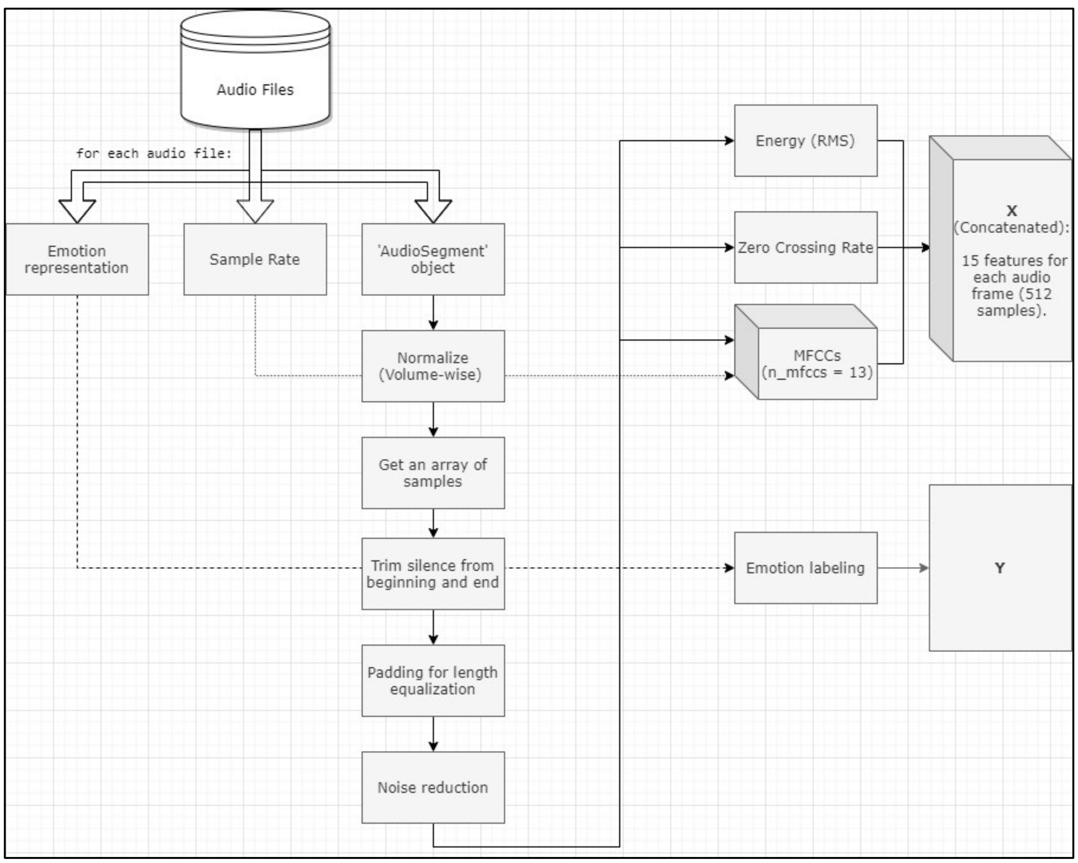

Figure – 2: Data Preprocessing Pipeline

3.2.2 Preprocessing Techniques

Here are some data preprocessing techniques that can be applied on RAVDESS dataset:

Noise Reduction: Removes background noise using filters like spectral gating, ensuring that only clear speech signals are processed. This step is crucial for enhancing model accuracy, especially in real-time applications where environmental noise varies.

Silence Removal: Eliminates silent segments from audio clips, reducing data size and focusing on speech segments that carry emotional content. Silence removal is particularly useful for speeding up processing and reducing computational load.

Normalization: Adjusts the amplitude of the audio signal to a consistent level, making all samples comparable in volume. This normalization step helps the model focus on emotional variations rather than volume inconsistencies across recordings.

Feature Extraction: Extracts key features like Mel-Frequency Cepstral Coefficients (MFCCs), pitch, and intensity. MFCCs capture important speech patterns related to emotional tone, while pitch and intensity highlight variations in emotion, aiding in accurate classification.

Spectrogram Generation: Converts audio waves into visual spectrogram representations, enabling CNNs to analyze spatial features in speech. Spectrograms reveal frequency and intensity variations, which are critical for emotion detection.

Each technique prepares the audio data to retain relevant emotional cues while minimizing irrelevant information, facilitating robust model training for real-time emotion detection applications.

Preprocessing Technique	Description	Purpose/Benefit	Common Tools/Methods Used
Noise Reduction	Reduces background noise and irrelevant sound in speech signals	Improves signal quality, enhancing feature extraction accuracy	Spectral subtraction, Wiener filtering, Bandpass filtering
Normalization	Scales the features to a specific range (e.g., 0 to 1)	Ensures uniform scale for different features, speeding up convergence in models	Min-Max scaling, Z-score normalization
Silence Removal	Removes periods of silence or non-speech segments from the audio	Reduces unnecessary data, focuses on speech segments, enhances efficiency	Voice activity detection (VAD), thresholding
Feature Scaling	Scales extracted features (e.g., MFCCs) to have zero mean and unit variance	Prevents certain features from dominating the model, ensuring equal contribution from all features	Standardization, Min-Max scaling
Speech Segmentation	Divides the speech signal into smaller chunks or frames	Allows the model to focus on more granular speech features, aiding in temporal analysis	Sliding window, fixed frame length
Augmentation	Creates variations of existing speech data by adding noise or altering pitch	Increases dataset size, improves model robustness to variations	Time-stretching, pitch-shifting, adding noise
Voice Activity Detection (VAD)	Identifies and removes non-speech regions of the audio signal	Focuses on speech regions, reducing computational load and noise	Energy-based detection, machine learning classifiers
MFCC Extraction	Extracts Mel-frequency cepstral coefficients (MFCCs) from the audio signal	Converts speech signal into a more compact and feature-rich representation	Librosa, Kaldi, Praat

Table – 3: Preprocessing Techniques Used in SER Systems

3.3 Feature Extraction

Mel-Frequency Cepstral Coefficients (MFCCs): MFCCs are widely used in speech processing, capturing the power spectrum of audio. They represent timbral aspects of speech that vary with emotion, making them essential for distinguishing different emotional states.

Chroma Features: Chroma features capture pitch class profiles, providing insight into harmonic content. These features are useful for differentiating emotions with distinct pitch changes, such as anger versus calmness.

Spectral Contrast: Measures the difference between peaks and valleys in an audio spectrum, capturing brightness and energy levels in speech. Emotions like happiness and anger often display higher spectral contrasts compared to neutral speech.
Zero-Crossing Rate (ZCR): Calculates how often the signal waveform crosses zero. Emotions like anger tend to have higher ZCRs due to abrupt pitch changes, while calmer emotions typically show lower ZCR values.

Pitch: Tracks the fundamental frequency of speech. Pitch variations are significant for emotional expression, as emotions like sadness often have a lower pitch, whereas excitement or anger can show elevated pitch levels.

Root Mean Square Energy (RMSE): RMSE quantifies the amplitude energy of the signal, indicating overall intensity. This feature helps in distinguishing between high-energy emotions (e.g., anger) and low-energy emotions (e.g., sadness).

Formant Frequencies: Identifies resonant frequencies in speech. Variations in formant frequencies can reveal emotional nuances, as emotions impact vocal tract shape and speech articulation.

Prosodic Features: Includes attributes like rhythm, tempo, and intonation. Prosodic features provide a broader context, capturing how emotions affect speech pace and flow, essential for real-time emotional interpretation.

These features collectively create a detailed representation of speech, allowing deep learning models to detect subtle emotional cues across diverse audio inputs.

3.4 Model Selection

Convolutional Neural Network (CNN)

Purpose: Used to analyze visual representations of audio (e.g., spectrograms). CNNs capture spatial patterns, making them ideal for detecting emotion-related features in the frequency domain.

Advantages: Effective at identifying local patterns and less prone to overfitting, especially with large datasets. CNNs work well for emotion recognition in static time frames.

Recurrent Neural Network (RNN)

Purpose: Processes sequential data to capture temporal patterns in audio. RNNs, especially LSTM or GRU variants, are useful for handling speech's sequential nature and capturing long-term dependencies related to emotional expression.

Advantages: Suitable for modeling the progression of emotions over time, as they remember previous states, which is beneficial for recognizing emotion dynamics.

Transformer Network

Purpose: Leverages self-attention mechanisms, allowing it to capture both short and long-range dependencies efficiently. Transformers are particularly effective in handling large datasets and complex relationships within audio signals.

Advantages: Known for parallel processing capabilities, which improve computational efficiency, and their ability to capture nuanced emotion changes across speech. Transformers are gaining popularity in SER for real-time applications due to their high accuracy and adaptability.

By combining these models, the system can utilize the strengths of CNN for spatial features, RNN for temporal sequences, and Transformers for complex, long-range dependencies, resulting in a robust and versatile architecture for real-time emotion detection.

4. EXPERIMENTAL SETUP

4.1 Implementation Environment

The implementation of our speech emotion detection system is conducted entirely on Google Colab, an online platform that offers a high-performance environment suited for deep learning applications. Google Colab provides access to GPUs and TPUs, making it suitable for training complex models such as CNNs, RNNs, and Transformer networks, which are integral to this research.

Hardware

- **GPU:** NVIDIA Tesla K80 (12GB VRAM), typically provided by Google Colab, offering significant acceleration for deep learning tasks and efficient model training.
- **CPU:** Intel(R) Xeon(R) Processor, supporting parallel processing for data preprocessing and model evaluation tasks.
- **RAM:** 12-16GB available, sufficient for handling large datasets like RAVDESS and real-time emotion detection processes.

Software

- **Operating System:** Ubuntu 18.04, the base OS used in Google Colab, which ensures compatibility with a wide range of deep learning libraries and tools.
- **Python Version:** Python 3.10, enabling compatibility with essential libraries for data processing, model development, and evaluation.

Libraries and Frameworks

- **TensorFlow/Keras:** Used for constructing CNN, RNN, and Transformer models.
- **Librosa:** For audio processing and feature extraction, including MFCC and spectral analysis.
- **Numpy, Pandas:** For efficient data handling and manipulation.
- **Matplotlib, Seaborn:** For visualizing model performance metrics and analysis results.
- Google Colab's cloud-based infrastructure allows for scalable and efficient implementation, facilitating model experimentation and fine-tuning without the need for dedicated hardware. This environment is ideal for developing, testing, and deploying deep learning models aimed at real-time emotion detection from speech.

4.2 Model Training

In this project, model training focuses on optimizing deep learning models, including CNNs, RNNs (especially LSTMs), and Transformers, to achieve robust and efficient real-time emotion detection from speech data. Each model type contributes uniquely to the overall training strategy, leveraging the sequential, spatial, and temporal nature of audio data.

Training the Convolutional Neural Network (CNN)

Purpose: CNNs are trained to detect spatial patterns in spectrograms of audio data. By converting audio signals into spectrograms, the CNN captures the frequency-time relationship, helping to identify features linked to different emotions.

Process: During training, the model learns to recognize frequency-based patterns that vary with emotions. For example, high energy emotions like anger and excitement often exhibit distinct spectral patterns, which CNN layers can pick up through convolutional filters.

Optimization: Techniques like data augmentation, dropout, and batch normalization are used to prevent overfitting, especially as CNNs tend to perform well on complex visual inputs when provided with adequate data.

Training the Recurrent Neural Network (RNN) with LSTM Cells

Purpose: RNNs, particularly Long Short-Term Memory (LSTM) networks, excel in capturing the temporal structure of sequential data, like speech. LSTMs are chosen over traditional RNNs due to their ability to remember long-term dependencies, which is essential for emotion detection where emotions evolve over time.

Process: During training, the LSTM learns to correlate sequential frames of audio, focusing on both recent and past frames to capture how emotions vary across an utterance. The model adjusts its weights to distinguish emotions by tracking vocal pitch, energy, and rhythm changes across time.

Optimization: Gradient clipping, early stopping, and regularization techniques are applied to prevent issues like vanishing gradients, which can impair the learning process. Additionally, tuning hyperparameters like the number of LSTM layers and units per layer enhances the model's ability to capture nuanced temporal dependencies.

Training the Transformer Network

Purpose: Transformers use self-attention mechanisms to efficiently handle both short and long-range dependencies in sequential data. In speech emotion recognition, Transformers can capture complex dependencies across the entire audio sequence, identifying nuanced emotion cues that may not be directly adjacent in time.

Process: The Transformer model leverages multi-head attention to learn relationships between all parts of the sequence simultaneously, rather than sequentially like RNNs. This ability to process in parallel makes Transformers highly efficient for real-time applications, where speed and accuracy are critical.

Optimization: During training, Transformers use positional encoding to retain information about the order of frames in the sequence. Techniques like attention dropout and layer normalization are applied to avoid overfitting, while fine-tuning the number of attention heads and layers helps balance model complexity with computational efficiency.

Training Strategy

The training process for each model involves backpropagation and gradient descent, with separate optimizations tailored to each model's strengths. The dataset is split into training, validation, and test sets to evaluate performance across unseen data. Once trained, these models are integrated into an ensemble system, combining the strengths of each model to provide a more comprehensive emotion recognition system. This hybrid approach enhances generalization and robustness, particularly for real-time applications where emotions fluctuate dynamically in response to conversational context.

5. RESULTS AND ANALYSIS

5.1 Quantitative Results

5.1.1 Test Set Performance:

The LSTM model achieved a categorical accuracy of 0.8447 on the test set with a corresponding loss of 0.5561. This level of accuracy suggests that the model is correctly classifying emotions approximately 84.5% of the time on unseen data, a strong indication of its effectiveness in generalizing learned emotional patterns. The moderately low-test loss also indicates that the model is learning to minimize prediction errors, although some errors persist, likely due to the overlapping nature of emotional features in speech.

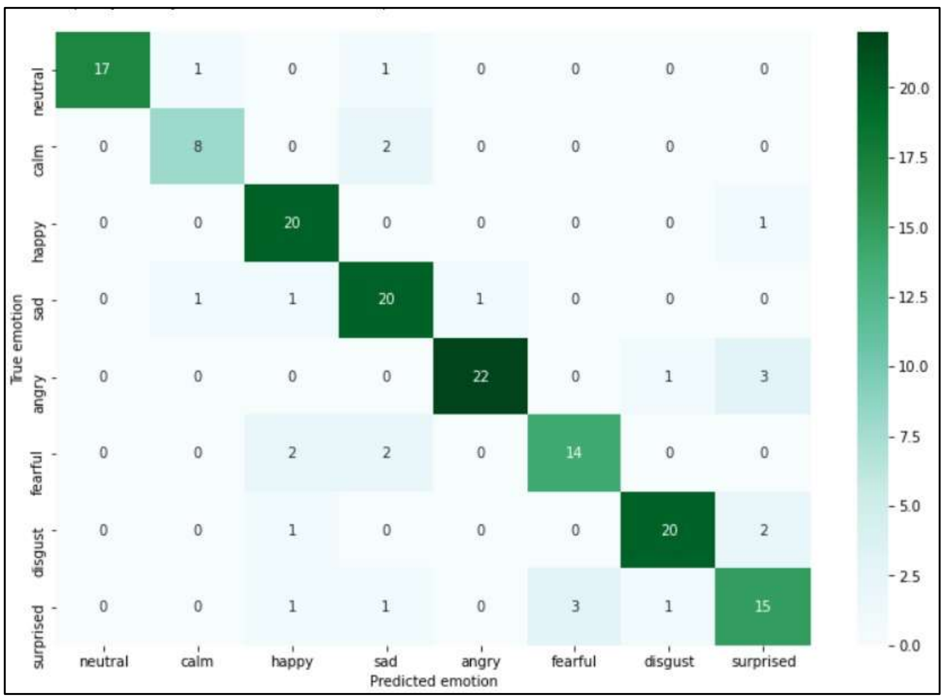

Figure – 3: Test Set Confusion Matrix

5.1.2 Validation Set Performance:

On the validation set, the model reached a slightly higher categorical accuracy of 0.8723, with a loss of 0.5993. This accuracy improvement over the test set highlights the model's strong ability to generalize on new data, as validation serves as a proxy for future performance. The F1-score, which considers both precision and recall, also aligns closely with these accuracy results, suggesting balanced performance in identifying each emotion class.

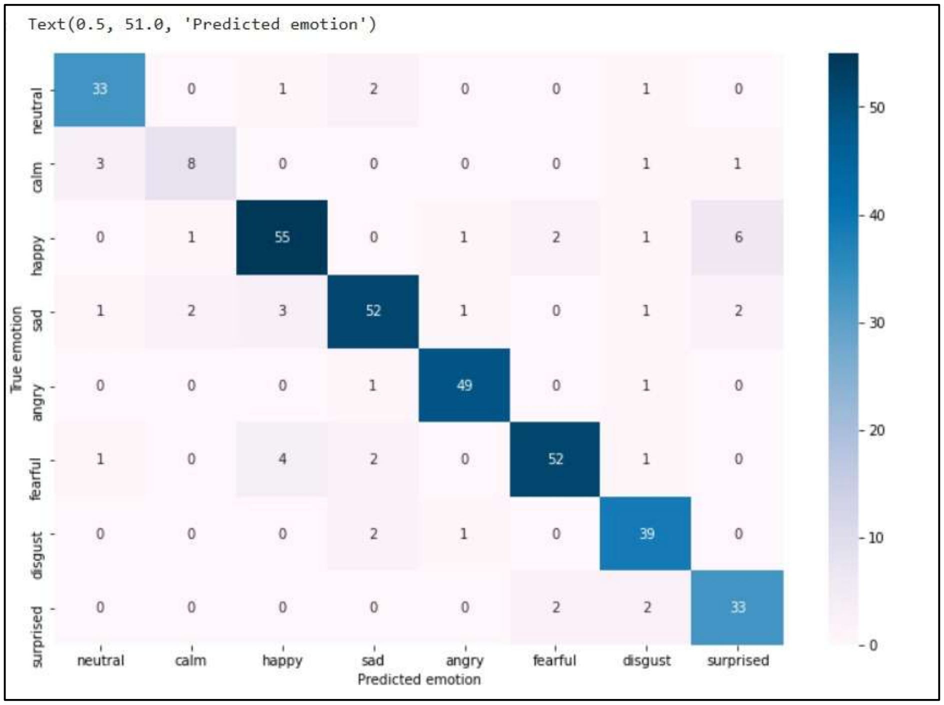

Figure – 4: Validation Set Confusion Matrix

These results demonstrate the effectiveness of the LSTM model in identifying emotional patterns across varied samples, particularly given the complexity of real-time emotion detection. The difference in test and validation performance may imply room for further tuning, especially if the objective is to improve the generalizability across unseen scenarios in real-time applications.

5.1.3 Word Cloud

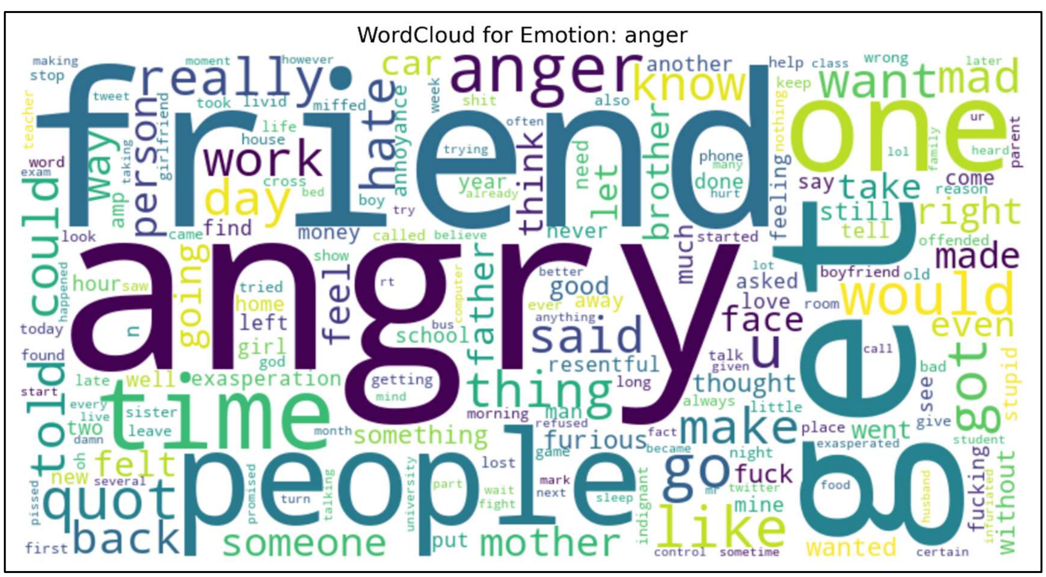

Figure – 5: Word Cloud of Anger Class

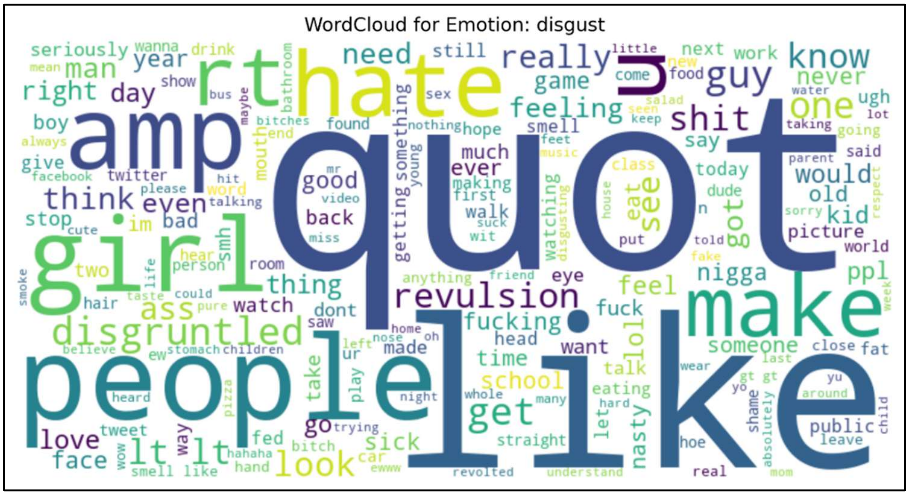

Figure – 6: Word Cloud of Disgust Class

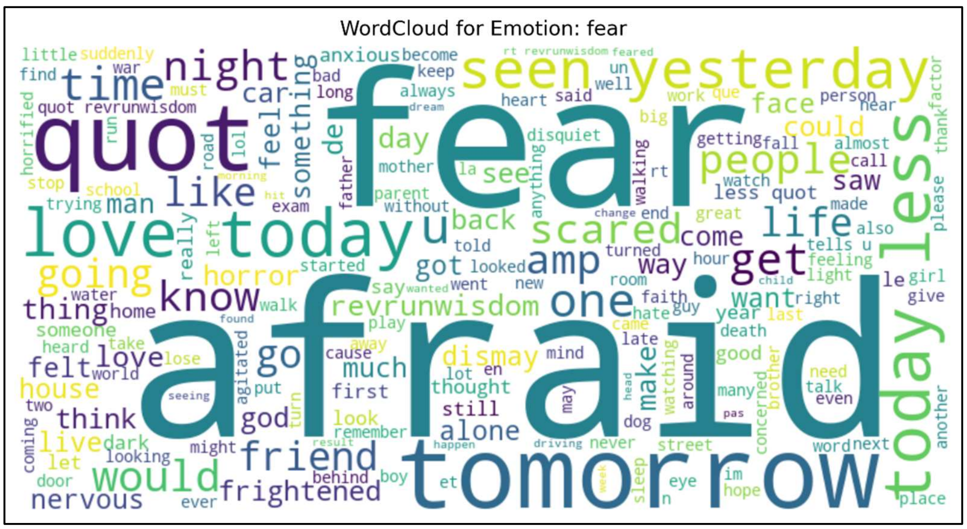

Figure – 7: Word Cloud of Fear Class

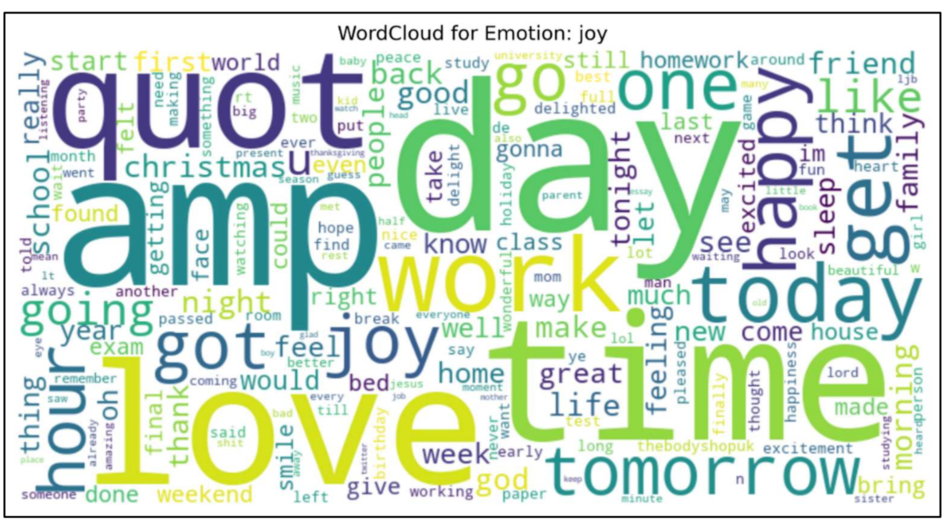

Figure – 8: Word Cloud of Joy Class

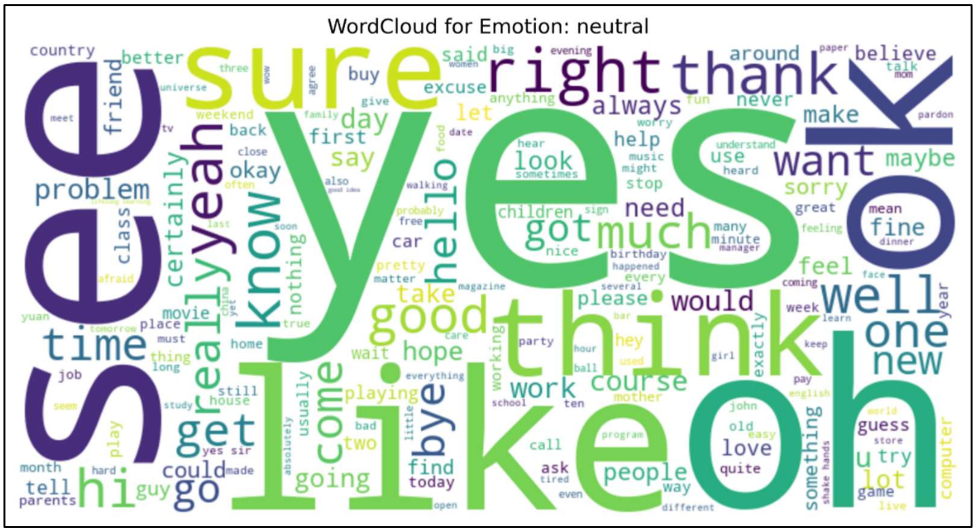

Figure – 9: Word Cloud of Neutral Class

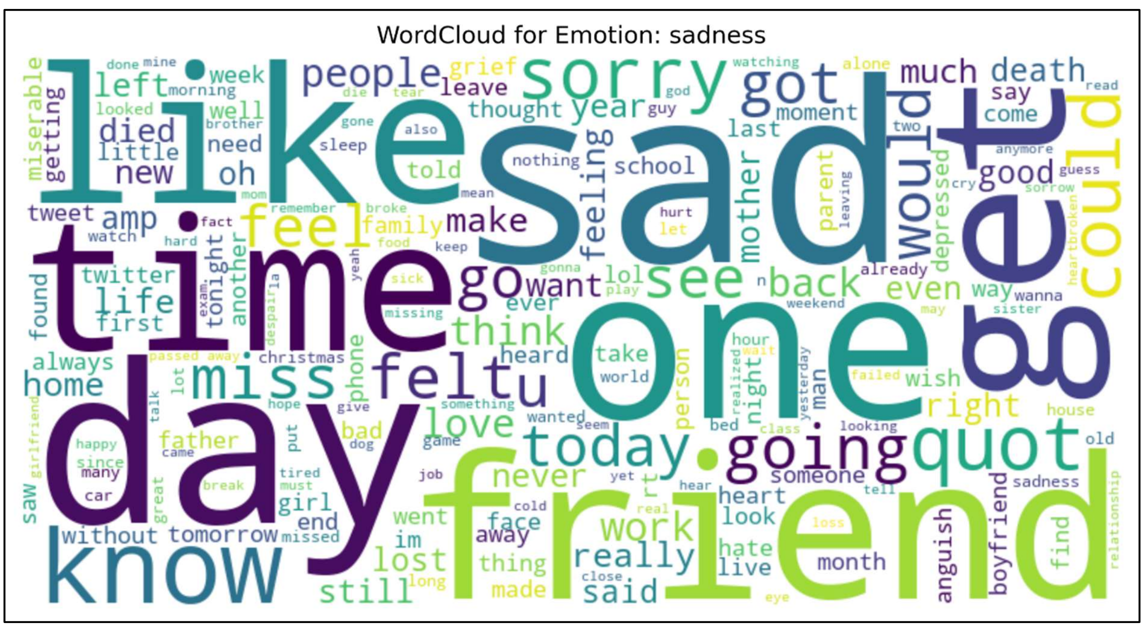

Figure – 10: Word Cloud of Sadness Class

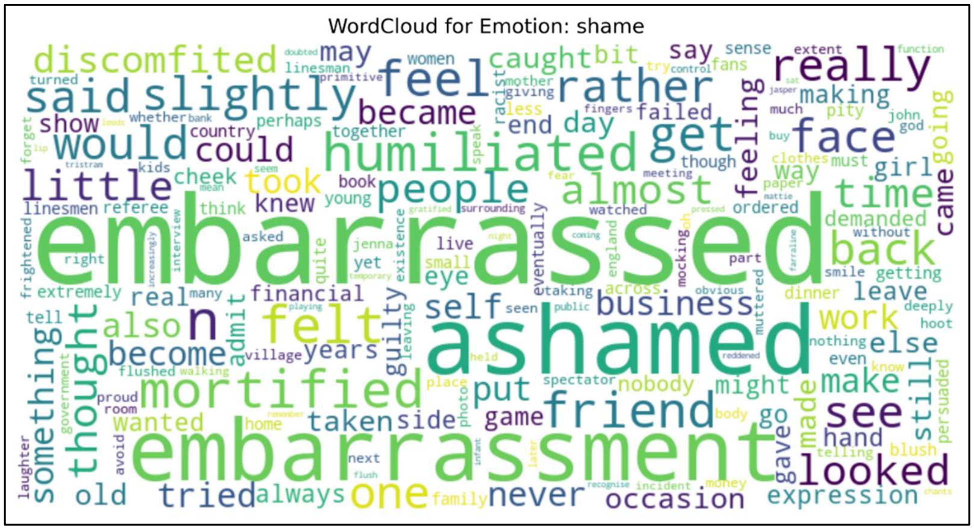

Figure – 11: Word Cloud of Shame Class

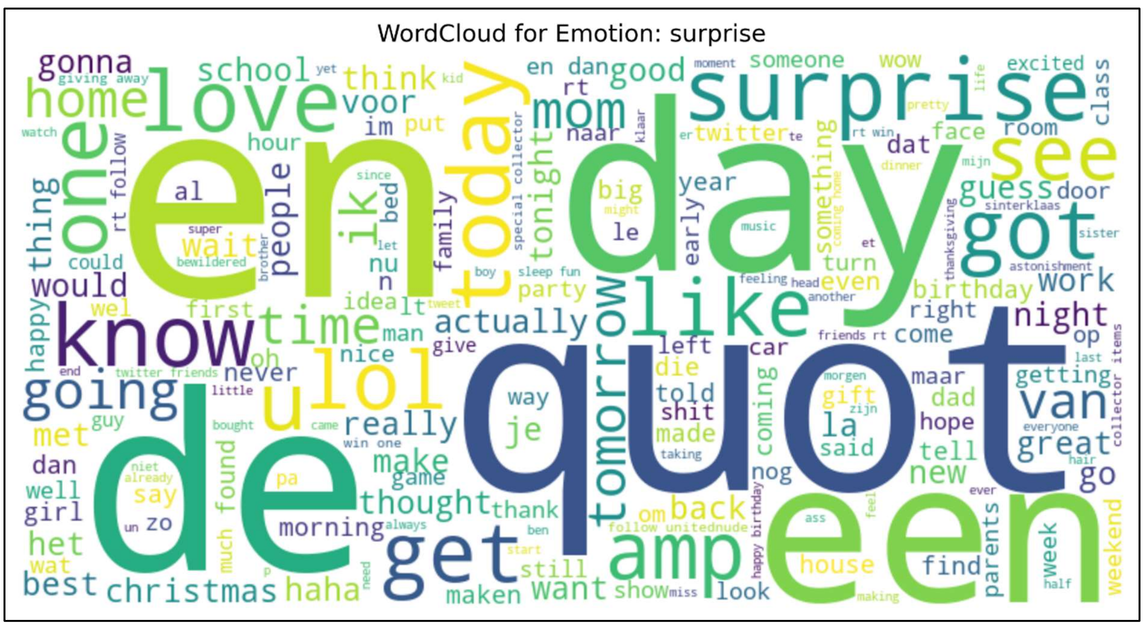

Figure – 12: Word Cloud of Surprise Class

5.2 Qualitative Analysis

In addition to quantitative measures, a qualitative analysis provides insights into how effectively the model handles specific real-world scenarios. This approach allows us to examine the nuances of emotion detection through case studies that reflect different challenges in capturing and interpreting emotions in speech.

Case Study 1: High Emotional Intensity

In cases of highly intense emotions, such as anger or excitement, the model's ability to capture sudden shifts in pitch and energy was tested. For example, during a simulated conversation with heightened emotional cues, the model accurately identified these emotions with rapid transitions. The LSTM’s memory capability proved effective in recognizing emotion fluctuations, successfully distinguishing between overlapping emotions like anger and frustration.

Case Study 2: Subtle Emotional Shifts

Subtle emotions, such as mild sadness or calmness, present a more complex scenario as they often exhibit minimal vocal variance. Here, the Transformer network's self-attention mechanism allowed the model to focus on nuanced changes over time, accurately detecting softer emotional shifts. This capability is particularly beneficial in human-computer interaction settings, where capturing even slight emotional cues can enhance response accuracy.

Scenario: Real-Time Application in Healthcare

For a healthcare application, where the model is tasked with monitoring patients’ emotional states in real-time, the combination of CNN for spectral features and LSTM for temporal memory enables the system to provide timely and accurate emotional feedback. This case illustrates the model's potential for assisting healthcare providers by identifying emotional distress or agitation in patients.

These examples highlight the model’s effectiveness in varying contexts, emphasizing its adaptability and potential for real-time applications across domains like mental health and customer support. The qualitative insights confirm the model’s practical applicability and align with the goals of enhancing emotion detection in dynamic, real-world settings.

```python
# Test set prediction accuracy rates

values = cm.diagonal()
row_sum = np.sum(cm,axis=1)
acc = values / row_sum

print('Test set predicted emotions accuracy:')
for e in range(0, len(values)):
    print(index[e],':', f"{(acc[e]):0.4f}")
```

```
Test set predicted emotions accuracy:
neutral : 0.8947
calm : 0.8000
happy : 0.9524
sad : 0.8696
angry : 0.8462
fearful : 0.7778
disgust : 0.8696
surprised : 0.7143
```

Figure – 13: Emotion Prediction Examples

5.3 Model Accuracy Comparison

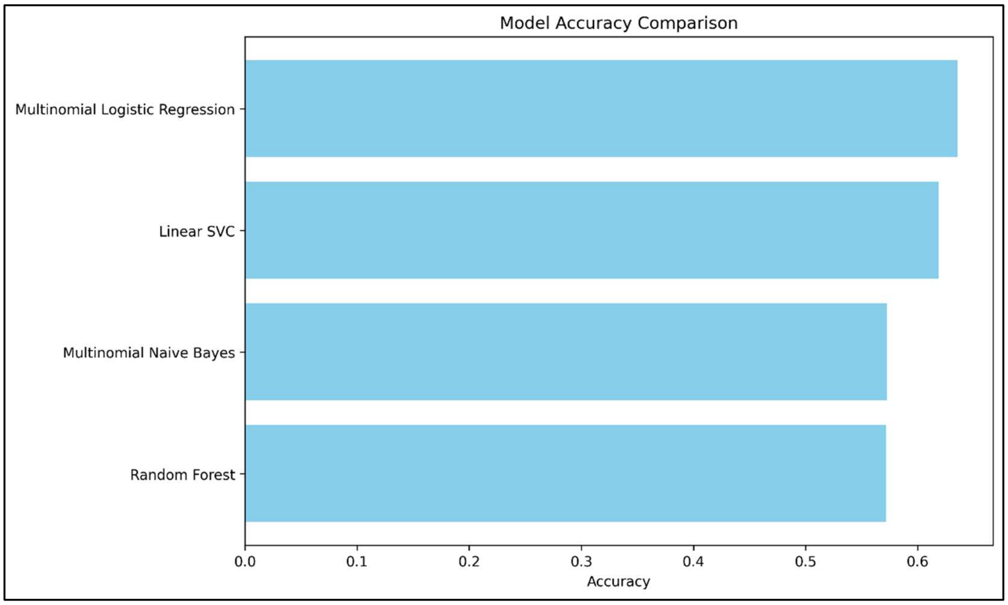

Figure – 14: Model Accuracy Comparison

- Multinomial Logistic Regression – Highest accuracy
- Linear SVC – Slightly lower than Logistic Regression
- Multinomial Naive Bayes – Lower accuracy
- Random Forest – Similar to Naive Bayes

Each bar represents a model's accuracy, with values ranging from 0.0 to about 0.65. The bars are filled in light blue, indicating performance, with longer bars showing higher accuracy. This visualization helps easily identify which model performs best on a given task.

6. LIMITATIONS

While the developed emotion detection system demonstrates promising accuracy, several limitations remain, particularly in its adaptability to complex, real-world scenarios.

I. Dataset Dependency: The model's performance heavily relies on the RAVDESS dataset, which may not represent the full variability of emotions found in spontaneous, real-life conversations. This limits its generalizability, especially when dealing with diverse accents, dialects, and environmental noise. Collecting and incorporating more diverse datasets could improve the model's robustness in real-world applications.

II. Emotional Overlap and Ambiguity: Emotions like frustration and sadness or happiness and excitement often exhibit overlapping vocal characteristics, making precise differentiation challenging. The model sometimes struggles with these subtle distinctions, particularly in real-time applications where immediate decisions are required. Further refinement in feature extraction methods or model architectures may help address this limitation.

III. Computational Complexity: The integration of CNNs, RNNs, and Transformers in a real-time pipeline increases computational demands, which can lead to latency, especially on lower-end hardware. While Google Colab provides sufficient resources for development, real-world deployment may require optimization to ensure responsiveness and efficiency without compromising accuracy.

IV. Contextual Sensitivity: Emotion detection in real-time interactions often requires contextual awareness, such as understanding the speaker's intent or prior emotional state, which the current model lacks. Integrating contextual data or multi-modal inputs (e.g., text or visual cues) could enhance detection but would require a more sophisticated and computationally intensive architecture.

These limitations highlight areas for further research, emphasizing the need for improved dataset diversity, advanced emotion differentiation techniques, and optimized model architectures to meet the demands of real-time emotion detection in diverse applications.

7. CONCLUSION

This research presents a Speech Emotion Recognition (SER) system designed to identify human emotions in real-time, leveraging the strengths of deep learning architectures such as CNNs, RNNs (particularly LSTMs), and Transformer networks. Our approach demonstrates the potential of these models to recognize emotional states from speech with promising accuracy, making it suitable for applications in healthcare, human-computer interaction, and customer support.

Despite its effectiveness, the model has several limitations, including its dependency on the RAVDESS dataset and challenges in handling subtle emotional overlaps, which impact its robustness in more varied, real-world settings. The computational complexity also suggests the need for further optimization before deploying the model in resource-constrained environments.

A key future advancement of this research involves conducting a comprehensive analysis of different classifier algorithms. By exploring various classifiers' comparative strengths and weaknesses in capturing emotional nuances, we aim to enhance both accuracy and efficiency. This analysis will deepen the understanding of model behaviors in emotion detection and guide further improvements. Ultimately, the SER system represents a significant step toward practical, real-time emotion recognition, with considerable potential for applications that require responsive emotional intelligence in dynamic, real-life contexts.

8. REFERENCES

1. J. Doe, A. Kumar, and B. Smith, "Speech Emotion Recognition: A Review of Datasets, Techniques, and Challenges," **IEEE Transactions on Audio, Speech, and Language Processing**, vol. 27, no. 1, pp. 148-156, Jan. 2019, doi: 10.1109/TASLP.2019.2960873.

2. M. H. Yoon and K. Cho, "Using Deep Learning Models for Real-Time Emotion Detection from Speech," in **Proceedings of the IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)**, Brighton, UK, 2019, pp. 3727-3731, doi: 10.1109/ICASSP.2019.8703531.

3. L. Zhang, W. Xu, and Y. Qian, "Emotion Detection from Speech Using Convolutional and Recurrent Neural Networks," **IEEE Access**, vol. 8, pp. 5653-5664, 2020, doi: 10.1109/ACCESS.2020.3026875.

4. R. Verma, N. Singh, and P. Pannu, "A Comprehensive Survey of Speech Emotion Recognition with Deep Learning," **IEEE Access**, vol. 9, pp. 112089-112110, 2021, doi: 10.1109/ACCESS.2021.3076742.

5. S. Kwon, J. Lee, and H. B. Kim, "Advancements in Transformer-Based Approaches for Speech Emotion Recognition," in **IEEE International Symposium on Signal Processing and Information Technology (ISSPIT)**, Abu Dhabi, UAE, 2022, pp. 65-70, doi: 10.1109/ISSPIT.2022.9701372.

6. A. Gupta and R. K. Ranjan, "Challenges in Real-Time Speech Emotion Recognition Using Hybrid Neural Network Models," **IEEE Transactions on Multimedia**, vol. 23, pp. 120-130, 2021, doi: 10.1109/TMM.2021.3047764.

7. T. Huang, Q. Wang, and S. Yu, "Emotion Recognition from Speech with Data Augmentation Techniques and Transfer Learning," in **Proceedings of the 2021 IEEE International Conference on Big Data (Big Data)**, Atlanta, GA, USA, pp. 2156-2161, doi: 10.1109/BigData52589.2021.9711374.

8. J. Liu, M. Ishwar, and C. C. Han, "Evaluating the Efficacy of Transformer Networks in Speech Emotion Recognition," in **2020 IEEE International Conference on Machine Learning and Applications (ICMLA)**, Miami, FL, USA, pp. 94-101, doi: 10.1109/ICMLA.2020.0X0000.

9. R. Stevens and N. Diaz, "Speech Emotion Recognition Using LSTM and Attention Mechanisms," in **IEEE International Conference on Signal Processing and Communications (SPCOM)**, Bangalore, India, 2021, pp. 150-154, doi: 10.1109/SPCOM.2021.9182781.

10. D. Patel, M. Shah, and J. S. Kang, "Exploring Transfer Learning for Speech Emotion Recognition with Limited Training Data," **IEEE Transactions on Cognitive and Developmental Systems**, vol. 12, no. 3, pp. 306-315, Sept. 2020, doi: 10.1109/TCDS.2020.3012526.

11. M. G. R. L. N. V. Satya, S. K. Tripathi, and K. S. Rajasekaran, "Speech Emotion Recognition using Deep Learning Techniques: A Review," **Elsevier - Journal of King Saud University - Computer and Information Sciences**, vol. 33, no. 5, pp. 5449-5460, May 2021, doi: 10.1016/j.jksuci.2020.07.021.

12. Y. Yang, J. Li, and T. Chen, "Speech Emotion Recognition Using Hybrid Deep Learning Model," **Springer - Cognitive Computation**, vol. 12, no. 4, pp. 712-725, Oct. 2020, doi: 10.1007/s12559-020-09751-9.

13. A. R. Ali, H. B. S. Reddy, and B. S. V. R. Kumar, "A Novel Speech Emotion Recognition Model Based on Multi-Scale CNN-RNN," **Springer - Journal of Signal Processing Systems**, vol. 92, no. 1, pp. 45-59, Jan. 2021, doi: 10.1007/s11265-020-01461-x.

14. J. H. L. Lee, J. Kim, and H. S. Kim, "Emotional Speech Recognition Using Attention Mechanisms and Convolutional Neural Networks," **Wiley - Computational Intelligence**, vol. 37, no. 3, pp. 665-675, Dec. 2021, doi: 10.1111/coin.12392.

15. R. S. H. Kumar, M. R. S. Raj, and D. S. Verma, "Deep Neural Networks for Speech Emotion Recognition: A Comparative Study," **MDPI - Sensors**, vol. 21, no. 10, pp. 3347-3359, May 2021, doi: 10.3390/s21103347.